

Clock Skew Scheduling: Formulations and Algorithms

Wai-Shing Luk

October 9, 2026

Abstract

Clock skew scheduling (CSS) assigns clock arrival times so that every setup- and hold-time constraint is met while a secondary objective is optimized. This survey rests on one idea: the timing constraints of a synchronous circuit form a system of difference constraints, feasible when the timing constraint graph has no negative cycle. Period minimization, slack maximization, and yield-driven scheduling reduce to parametric potential problems solved by Bellman–Ford, Lawler, Howard, Karp, and the ellipsoid method. We extend the core with delay padding, statistical methods, GEV delay models, and clock-tree co-optimization, closing with idealizations, common misconceptions, and open problems.

1 Introduction

Synchronous digital design rests on a global clock that must reach every sequential element—every flip-flop (FF) and latch—at the right instant. Because the clock distribution network has finite, unequal delays, the active clock edge does not arrive at all registers simultaneously. The difference in arrival times is called *clock skew*. For a data path from register i to register j , with arrival times u_i and u_j , the (useful) skew is

$$y_{ij} = u_i - u_j.$$

Two facts make skew a first-class design variable. First, the clock is *periodic*, so a common shift $u_i \leftarrow u_i + c$ changes nothing: only the *differences* y_{ij} matter. Second, the setup- and hold-time constraints that guarantee correct data capture bound each skew from above and below, so each data path defines an interval of admissible skews. Clock skew scheduling (CSS) is the problem of choosing the arrival times—equivalently the skews—so that every timing constraint is met and some secondary objective is optimized.

Historically, designers targeted *zero skew* ($u_i = u_j$), which is easy to reason about: a single worst-case clock period can be verified by ordinary static timing

analysis (STA). Zero skew, however, throws away a valuable degree of freedom. *Useful-skew* design [1], [2], [3] deliberately introduces nonzero skew to borrow slack from non-critical paths and lend it to critical ones. The result is either a higher clock frequency, a larger timing margin, or—under process variations—a higher parametric yield.

1.1 Why clock skew scheduling?

The value of a skew schedule is easily underestimated, because a schedule is a *plan* rather than a physical artifact and it is therefore tempting to dismiss it as an academic exercise. The practical considerations below argue otherwise.

The flow order must respect skew scheduling. Scheduling requires accurate clock and data delays, which presuppose that global and clock-tree routing have already been performed. The conventional flow, however, treats clock-tree synthesis (CTS) as the most timing-critical step and performs it *before* routing the signal nets. A flow that is aware of useful skew would invert that priority: run CTS first, as if the signal wires were not yet fixed, and then re-route only the few nets whose geometry the clock tree displaces. This “CTS first, then re-route” ordering—rarely discussed—determines how much of the available scheduling freedom can actually be realized.

Placement is the true lever. If the placement is poor, no algorithm inside CTS can balance the clock tree, because the imbalance is a property of the register locations, not of the clock tree. Placement must therefore anticipate the scheduler. The consequence is that useful-skew design is less a single algorithm than a *tool-chain contract*: placement, routing, and CTS must each be aware of, and cooperate with, the skew schedule.

A timetable, not a single departure time. The clearest intuition is a railway timetable. Trains do not all depart at the same instant, nor are they required to arrive simultaneously at every station; departures and arrivals are deliberately staggered to maximize the throughput of the network. Zero-skew clocking

is the opposite policy—every register is clocked at the same instant—and it forgoes precisely the freedom that a timetable exploits. Useful skew is the timetable of the clock, and the asymmetry of current investment—large sums directed at lithography while the comparatively inexpensive software freedom of skew scheduling is left unexploited—becomes harder to justify as devices shrink and variation grows.

Why the formulation survives increasing complexity. The mathematical core of the problem is fixed at the outset and does not change as the model grows. The governing invariant is that only *relative* arrival times matter: shifting every clock by the same amount is physically meaningless, just as it is immaterial whether a schedule starts on Christmas or on New Year’s Day. This invariant is exactly the hypothesis under which the constraints form a system of difference constraints, so the negative-cycle feasibility test and the parametric machinery of Sections 4–8 survive the statistical refinements almost unchanged. The scalar bisection of the single-parameter case generalizes to a Newton- or ellipsoid-style iteration when several parameters are present (Section 4.4); entirely new machinery is required only if a constraint is placed on *absolute* time, a situation that current practice does not produce. This stability is the strongest reason to study the problem in the abstract form developed below.

1.2 The central idea

This survey is organized around a single unifying observation:

The timing constraints of a synchronous circuit form a *system of difference constraints*,

$$u_i - u_j \leq w_e \quad \text{for every arc } e = (j, i),$$

and such a system is feasible if and only if the associated *timing constraint graph* (TCG) contains no *negative cycle*.

Consequently, the natural formulation places clock skew scheduling within *network optimization*. Let A be the node–edge incidence matrix of the TCG, u the vector of arrival times, and $y = Au$ the vector of skews. Then the timing problem is a *feasible potential problem*:

$$\underline{w} \leq y = Au \leq \overline{w},$$

and its optimization variants—minimum clock period, maximum slack, maximum yield—are all *parametric potential problems* of the form

$$\max\{\beta \in \mathbb{R} : y \leq d(\beta), \quad Au = y\},$$

where $d(\beta)$ is a monotone decreasing function of the scalar parameter β . When $d(\beta) = m - s\beta$ is affine with nonnegative denominators s , this is the *minimum cost-to-time ratio cycle* problem; when s is a constant it further reduces to the *minimum mean cycle* problem. Both admit fast combinatorial algorithms—Howard’s policy iteration, Karp’s characterization, and Lawler’s parametric search—and the negative cycle returned by these algorithms is not merely a certificate of infeasibility but also the *most critical* set of paths, which is exactly the information a designer needs.

1.3 Scope and organization

The remainder of the article is organized as follows. Section 2 fixes notation, derives the setup- and hold-time constraints, and constructs the timing constraint graph. Section 3 develops the network-potential formulation and the negative-cycle feasibility certificate. Section 4 treats clock-period minimization and slack maximization as parametric shortest-path problems, and extends them to multiple parameters via the ellipsoid method. Section 5 generalizes the parametric formulation to convex and quasi-convex network problems, covering min-cost flow, cycle cancellation, line search, linear-fractional and statistical costs, and the barrier method. Section 6 covers delay padding and its physical configurations. Section 7 addresses yield-driven scheduling under process variations, including the EVEN, PROP, C-PROP, and FP-PROP methods. Section 8 extends the model to non-Gaussian, heavy-tailed delay distributions via the generalized extreme value (GEV) distribution. Section 9 compares the underlying algorithms and their complexities. Section 10 links the schedule to clock-tree synthesis, and Section 11 lists open problems.

This article is a *survey*: it organizes and explains known formulations and algorithms rather than reporting new experiments, and the quantitative speedups quoted below are those reported in the cited literature. Where the abstract model diverges from physical design—clock-tree power and area, discrete cell sizes, spatial correlation, and hierarchical timing—the discrepancy is stated explicitly and revisited in Section 11.

2 Preliminaries: Timing Constraints and the Timing Constraint Graph

2.1 Local data paths and clock skew

Consider a data path that starts at the output of register i , passes through a block of combinational logic, and ends at the data input of register j (Figure 1). The data signal traverses the logic in a delay that ranges between a minimum d_{ij} and a maximum D_{ij} over the relevant operating conditions. The two registers are clocked by edges whose arrival times at i and j are u_i and u_j ; the skew seen by the path is $y_{ij} = u_i - u_j$.




Figure 1: A local data path from register R_i to register R_j . The data traverses the combinational logic in a delay between d_{ij} and D_{ij} ; only the skew $y_{ij} = u_i - u_j$ of the two clock arrival times matters.




Figure 2: Clock waveforms. The source clock CLK_src and the clocks arriving at registers i and f are delayed by t_i and t_f ; their difference $t_i - t_f$ is the skew.

Figure 2 shows the corresponding clock waveforms: the arrival times are simply the delays of the clock edges at the two registers, and the skew is their difference.

The circuit as a whole is abstracted as a directed graph in which vertices are registers and edges are data paths. A small example, together with its timing

constraint graph, is shown in Figure 3.

2.2 Setup- and hold-time constraints

Correct capture requires two inequalities. The *setup-time* constraint says that data launched by i must reach j early enough before the capturing edge:

$$y_{ij} \leq T_{\text{CP}} - D_{ij} - T_{\text{setup}}. \quad (1)$$

Here T_{CP} is the clock period and T_{setup} the setup time of register j . A violation of (1) is a *cycle-time violation*, also called *zero clocking*. The *hold-time* constraint says that new data must not arrive at j so early that it races past the previous capture:

$$y_{ij} \geq T_{\text{hold}} - d_{ij}. \quad (2)$$

A violation of (2) is a *race condition*, also called *double clocking*. Combining the two, the admissible skew of path $i \rightarrow j$ lies in the *feasible skew region* (FSR)

$$\begin{aligned} \underline{w}_{ij} &\leq y_{ij} \leq \bar{w}_{ij}, \\ \underline{w}_{ij} &= T_{\text{hold}} - d_{ij}, \\ \bar{w}_{ij} &= T_{\text{CP}} - D_{ij} - T_{\text{setup}}. \end{aligned} \quad (3)$$

Because $y_{ij} = u_i - u_j$, the constraints (3) are exactly a system of *difference constraints* on the arrival times u .

2.3 The timing constraint graph

The FSR (3) is represented by a directed graph $G(V, E)$, the *timing constraint graph*, one vertex per register:

- each *setup* constraint contributes an *s*-edge from j to i with weight $\bar{w}_{ij} = T_{\text{CP}} - D_{ij} - T_{\text{setup}}$;
- each *hold* constraint contributes an *h*-edge from i to j with weight $-\underline{w}_{ij} = d_{ij} - T_{\text{hold}}$.

Indeed, the *s*-edge $j \rightarrow i$ encodes $u_i \leq u_j + \bar{w}_{ij}$, i.e., (1); the *h*-edge $i \rightarrow j$ encodes $u_j \leq u_i - \underline{w}_{ij}$, i.e., (2). A vertex assignment u satisfying every arc is precisely a *feasible potential*, and we have the fundamental equivalence:

The timing constraints are satisfiable by some clock schedule if and only if the timing constraint graph has no cycle of negative total weight.

Summing the constraints around a cycle C eliminates the arrival times, because $\sum_{e \in C} (u_{\text{head}} - u_{\text{tail}}) = 0$;

hence a cycle is feasible only if its weight is nonnegative. A negative cycle therefore certifies that *no* skew assignment can satisfy the corresponding constraints, and the only remedy is a circuit-level change—delay padding or logic restructuring (Section 6).

2.4 The zero-skew baseline and reported slacks

It is often convenient to start from a zero-skew analysis. If the clock is balanced so that $y_{ij} = 0$ for all paths, STA reports a *setup slack* S_{ij} and a *hold slack* H_{ij} for each path. Under useful skew, the constraints become

$$y_{ij} = u_i - u_j \leq S_{ij}, \quad -y_{ij} = u_j - u_i \leq H_{ij}.$$

Since $S_{ij} = T_{\text{CP}} - D_{ij} - T_{\text{setup}}$ is exactly the upper skew bound, the reported slacks are directly the FSR bounds relative to the zero-skew solution. This is why a modern STA engine that reports only slacks is already sufficient as a front end: the slacks *are* the edge weights of the TCG.

3 The Network-Potential Formulation

3.1 Flow, potential, and duality

The TCG is a *network* in the sense of discrete calculus. Its node–edge incidence matrix A maps node potentials to edge tensions, $y = Au$, while its transpose A^T is the boundary operator that maps edge flows to nodal divergences. A *flow* x satisfies conservation $A^T x = 0$ (a circulation); a *tension* y satisfies $\sum_{e \in C} y_e = 0$ around every cycle. Tellegen’s theorem, $x^T y = x^T A u = (A^T x)^T u = 0$, states that flow and tension are bi-orthogonal—the two faces of the same network (Figure 4).

3.2 The feasible potential problem

In this language the timing problem is the *feasible potential problem*

$$\text{find } u \quad \text{such that} \quad \underline{w} \leq y \leq \overline{w}, \quad Au = y. \quad (4)$$

Feasibility of (4) has a clean combinatorial characterization. Let $d^+(P) = \sum_{e \in P} d_e^+$ and $d^-(P) = \sum_{e \in P} d_e^-$ denote the *upper span* and *lower span* of a cycle P , where d_e^+ and d_e^- are the upper and lower bounds on edge e . Because $\tau^T y = \tau^T A u = 0$ for the

indicator vector τ of any cycle, the constraints imply $d^-(P) \leq 0 \leq d^+(P)$ for every cycle P . Conversely, these cycle conditions are sufficient, and the “if” direction is constructive: it is exactly the shortest-path computation that produces u . In particular, when only upper bounds are present after adding reverse edges, the condition simplifies to $d^+(P) \geq 0$ for all cycles, and an infeasible instance returns a *negative cycle*.

3.3 Bellman–Ford and the case for lazy evaluation

The classical instrument for this task is the Bellman–Ford algorithm [4], [5] (Algorithm 1). It computes shortest-path potentials and, in a final pass, either certifies feasibility or exhibits a negative cycle. Its simplicity is appealing, but it has practical drawbacks that matter in timing closure:

1. it is designed to find shortest paths, with negative-cycle detection only a by-product;
2. it detects a negative cycle only at the *end* of the computation;
3. it requires all edge weights w to be known up front.

In an interactive timing-closure loop, the recommendation is therefore *lazy evaluation*: run the analysis on demand, vertex by vertex, and **stop as soon as a negative cycle is found**. The early-exit variant turns the algorithm from a global pass into a targeted repair tool.


4 Clock Skew Scheduling: Feasibility and Period Minimization

4.1 Minimizing the clock period

A traditional objective is to find the smallest clock period T_{CP} for which a feasible schedule exists. This is a linear program,

$$\begin{aligned} &\text{minimize} && T_{\text{CP}} \\ &\text{subject to} && \underline{w}_{ij} \leq u_i - u_j \leq T_{\text{CP}} - D_{ij} - T_{\text{setup}}, \\ & && \forall i \rightarrow j, \end{aligned} \quad (5)$$

whose constraints separate into a *hold part* independent of T_{CP} and a *setup part* that shrinks linearly as T_{CP} decreases. For a fixed T_{CP} , feasibility is decided by a single negative-cycle test on the TCG. Writing



(a) Example circuit and its graph abstraction.

(b) Its timing constraint graph.

Figure 3: A circuit and its timing constraint graph. Solid arcs are setup constraints (s -edges); dashed arcs are hold constraints (h -edges). The weights are $T_{CP} - 4$ and 1.5, etc.

Algorithm 1: Negative-cycle feasibility test (Bellman–Ford).

Input: TCG $G = (V, E)$ with weights w_e

Output: a feasible potential u , or a negative cycle

```

1 foreach  $v \in V$  do
2    $u_v \leftarrow 0$ ;  $\pi_v \leftarrow \text{nil}$ 
3 for  $k \leftarrow 1$  to  $|V|$  do
4   foreach  $(p, q) \in E$  do
5     if  $u_q > u_p + w_{pq}$  then
6        $u_q \leftarrow u_p + w_{pq}$ ;  $\pi_q \leftarrow p$ ;
7 foreach  $(p, q) \in E$  do
8   if  $u_q > u_p + w_{pq}$  then
9     return the negative cycle obtained by tracing  $\pi$ ;
10 return  $u$ 

```




Figure 4: A flow (circulation) x and a tension (potential difference) y are dual objects on the same network.

the cycle condition explicitly, a cycle C containing $k \geq 1$ setup edges is nonnegative exactly when

$$T_{CP} \geq \frac{\sum_{e \in C^s} (D_e + T_{\text{setup}}) - \sum_{e \in C^h} (d_e - T_{\text{hold}})}{k}.$$

The smallest feasible period is therefore a *maximum cycle ratio*: it is obtained by maximizing a ratio of a “cost” sum to a “time” sum over cycles. This is the minimum-cost-to-time-ratio cycle problem, and when the denominator is constant across edges it becomes the minimum mean cycle problem.

4.2 The parametric potential problem

All of the objectives above are instances of one *parametric potential problem*:

$$\max\{\beta : y \leq d(\beta), Au = y\}, \quad (6)$$

$d(\beta)$ monotone decreasing.

If $d(\beta) = m - s\beta$ is affine with $s \geq 0$, (6) is the minimum cost-to-time ratio cycle problem; if s is a constant, it is the minimum mean cycle problem. If $d(\beta)$ is nonlinear, (6) still makes sense but the linear combinatorial machinery must be replaced by binary search or a convex solver. A useful sufficient condition is the following.

If $g(\beta)$ and every $f_e(\beta)$ are monotone decreasing, the parametric problem $\max\{g(\beta) : u_i - u_j \leq f_e(\beta)\}$ has a unique solution.

4.3 Cycle-based algorithms

Two families of algorithms solve (6) directly.

Binary search (Lawler) [6]. Maintain an interval $[\beta_{\min}, \beta_{\max}]$ bracketing the optimum. For a trial β , test the TCG for a negative cycle; if one exists, the trial is too optimistic and the upper bound moves

down, otherwise the lower bound moves up. Each test costs one Bellman–Ford pass, and the number of tests is logarithmic in the inverse tolerance. Lawler’s method is simple and robust but can be slow.

Cycle cancellation (Howard) [7]. Howard’s policy iteration maintains a schedule (equivalently, a set of shortest-path trees) and repeatedly finds a cycle along which the ratio can be improved, “zeroing out” that cycle with the smallest possible effort. It converges very fast in practice and, as a by-product, returns the sequence of critical cycles in decreasing order of criticality. An improved variant introduces binary-search steps to accelerate convergence; hybrid methods combine the strengths of both.

4.4 Multi-parameter problems and the ellipsoid method

Every formulation so far optimizes a *single* scalar parameter—the clock period, a uniform slack, or one yield figure—and is therefore solvable by the combinatorial algorithms above. A real design, however, is often governed by *several* parameters at once: the clock period together with the timing slack, several yield targets, or the periods of multiple clock domains. The general problem is

$$\max\{g(\beta) : t_i - t_j \leq f_{ij}(\beta) \quad \forall (i, j) \in E\}, \quad (7)$$

$\beta \in \mathbb{R}^p,$

where g and the f_{ij} are monotone in each component of β (and, for the convex extension, convex) but need not be affine. For $p = 1$ this is the parametric potential problem (6), solved by the minimum-ratio-cycle algorithms. For $p > 1$, parametric shortest-path theory no longer applies, and one must turn to a general convex optimization method.

The reason the problem stays tractable is that a *negative-cycle test is a separation oracle*. Given a trial β , either the timing constraint graph $G(\beta)$ is feasible, or it contains a negative cycle C ; in the latter case

$$\sum_{e \in C} f_e(\beta) < 0$$

is a violated constraint that separates β from the feasible set. An oracle therefore needs only *one* violated constraint per iteration and never requires all $|E|$ constraints—indeed, all path delays—to be known in advance. This is the “lazy evaluation” principle of Section 3 lifted from feasibility to optimization.

The *ellipsoid method* [8], [9] is the natural engine for such an oracle: it maintains an ellipsoid known to contain the optimum, and each oracle call supplies a separating half-space through the current center, after which the ellipsoid is shrunk to the appropriate side until its volume (and hence the optimality gap) drops below a tolerance. Because it needs only separation, it tolerates non-differentiable objectives and nonlinear constraints; conceptually it is the multi-dimensional generalization of the bisection strategy used by Lawler’s algorithm. Its update is

$$\beta_{k+1} = \beta_k - \frac{1}{p+1} \frac{A_k g_k}{\sqrt{g_k^\top A_k g_k}},$$

where A_k defines the current ellipsoid and g_k is the separating gradient; deep, central, and parallel-cut refinements accelerate convergence in practice.

This is exactly the architecture implemented by the `netoptim` package, whose `NetworkOracle` wraps Howard’s negative-cycle finder and, on detecting a negative cycle, returns a cutting plane—its gradient is the negative sum of the edge sub-gradients and its intercept the negative sum of the edge weights—which is consumed by the ellipsoid and cutting-plane routines of `ellaigo`. The resulting solver handles both linear and nonlinear convex multi-parameter clock skew scheduling; empirically it is reported to be up to an order of magnitude faster than a general linear-programming solver on the linear problem, and several hundred times faster than a general convex-programming solver on the nonlinear problem [10].

4.5 Slack maximization

When no timing violation exists, the natural goal is to make the circuit as robust as possible by *maximizing the minimum slack*. The EVEN method (Section 7) is the simplest instance: maximize β subject to $u_i - u_j \leq \mu_{ij} - \beta$, which is a minimum mean cycle problem. The optimal β is the amount of uniform safety margin that can be inserted into every constraint simultaneously. The *minimum balancing* (MB) algorithm [11] realizes this schedule constructively: it finds the most critical cycle, distributes the slack evenly along it, contracts it to a super-vertex, and repeats until a single vertex remains. The contraction order is itself valuable, because it records the *hierarchy of criticality* (Figure 5).




Figure 5: The criticality hierarchy produced by minimum balancing. Each internal node merges a critical cycle; deeper subtrees are less critical.

5 Convex and Quasi-Convex Network Problems

5.1 From a scalar parameter to a quasi-convex objective

Sections 4 and 4.4 reduced every scheduling variant to a parametric potential problem—(6) and (7)—in which the right-hand side $d(\beta)$ of a *scalar* parameter is monotone. Monotonicity, however, is a convenience rather than a requirement: the negative-cycle test that decides feasibility of $y \leq d(z)$, $Au = y$ works for any right-hand side d . The convex-optimization viewpoint therefore enlarges the admissible objectives from monotone functions of a scalar to *convex* and *quasi-convex* functions of a vector z , and it supplies the separation oracle—a negative cycle—through which the enlarged problems are solved.

The bridge from a quasi-convex objective to a convex feasibility test is the *level-set* (sublevel-set) characterization. A function f is quasi-convex exactly when every sublevel set $\{z : f(z) \leq \gamma\}$ is convex; equivalently, when there is a family of functions Φ_γ such that

- $\Phi_\gamma(z)$ is convex in z for every fixed γ ,
- $\Phi_\gamma(z)$ is non-increasing in γ for every fixed z , and
- $f(z) \leq \gamma$ if and only if $\Phi_\gamma(z) \leq 0$.

For a ratio of a convex function to a positive concave function, $f(z) = p(z)/q(z)$ with p convex, q concave, $p \geq 0$, and $q > 0$, the choice

$$\Phi_\gamma(z) = p(z) - \gamma q(z)$$

has all three properties: it is convex in z for each fixed γ , and its 0-sublevel set is exactly the γ -sublevel set of f . Linear-fractional functions, distance ratios, and the minimum cost-to-time ratio cycle of Section 4 are all of this form.

Minimizing a quasi-convex f therefore reduces to a sequence of convex feasibility problems parametrized

by γ ,

$$\begin{aligned} & \text{find} && z, \\ & \text{subject to} && \Phi_\gamma(z) \leq 0, \\ & && y \leq d(z), \quad Au = y, \end{aligned} \quad (8)$$

in which the subproblem is feasible if and only if $\gamma \geq f^*$ and infeasible otherwise; a bisection on γ converges to the optimum. Each subproblem is convex and its separation oracle is the negative-cycle test applied to the graph weighted by d : one violated constraint per iteration suffices, which is exactly the lazy-evaluation principle of Section 4.4 lifted from a scalar parameter to a vector.

A useful special case is *monotone min-max*,

$$\min_y \left\{ \max_{ij} f_{ij}(y_{ij}) : Au = y \right\},$$

with each f_{ij} non-decreasing. Introducing a scalar β and inverting the f_{ij} reconstitutes a monotone parametric potential problem,

$$\max \{ \beta : y_{ij} \leq f_{ij}^{-1}(\beta) \quad \forall (i, j), \quad Au = y \},$$

because each inverse is non-decreasing in β . The max-min yield formulation of Section 7 is precisely of this type, with f_{ij} the quantile function of a path-delay distribution; the framework below is what makes its probabilistic and nonlinear variants solvable by the same negative-cycle machinery.

5.2 Min-cost flow and cycle cancellation

The same primitive drives *min-cost flow*. The linear problem is

$$\begin{aligned} & \text{minimize} && d^\top x + p \\ & \text{subject to} && c^- \leq x \leq c^+, \\ & && A^\top x = b, \quad b(V) = 0, \end{aligned} \quad (9)$$

where x is an arc-flow vector, $A^\top$ the incidence matrix of G , c^- and c^+ (possibly infinite) capacity bounds, and b a balanced supply vector. Its classical algorithms fall into two families. *Augmenting-path* methods start from an infeasible flow and inject the smallest possible amount along an augmenting path, preserving infeasibility until no flow remains to inject. *Cycle-cancelling* methods [6], [12] start from a feasible flow x_0 and repeatedly set $x_1 = x_0 + \alpha \Delta x$, where Δx is the indicator of a negative cycle of the residual graph and $\alpha > 0$.

The second family is a descent method in disguise. Because Δx is a circulation, $A^\top \Delta x = 0$, a negative cycle with respect to the cost d satisfies $d^\top \Delta x < 0$, so moving along it decreases the cost. The step is limited by the residual capacities,

$$\alpha_1 = \min_{ij} \{ c_{ij}^+ - x_{ij}^0 : \Delta x_{ij} > 0 \}, \quad \alpha_2 = \min_{ij} \{ x_{ij}^0 - c_{ij}^- : \Delta x_{ij} < 0 \},$$

with $\alpha_{\text{lin}} = \min\{\alpha_1, \alpha_2\}$; if $\alpha_{\text{lin}} = +\infty$ the problem is unbounded. Choosing a *minimum mean* negative cycle rather than an arbitrary one accelerates convergence. Cycle cancellation is thus the network specialization of a generic descent method: choose a descent direction, choose a step size, update, and repeat.

This is the structure the clock-skew solvers exploit. The feasible-potential problem $\underline{w} \leq Au \leq \bar{w}$ is dual to a min-cost flow of the form (9), and the negative cycle returned by Howard's algorithm (Section 4) is simultaneously the infeasibility certificate and the descent direction. Delay padding (Section 6) is the primal min-cost-flow instance, whose dual variable x measures how the padding budget is spent.

5.3 Convex costs and line search

When the cost is a convex nonlinear function $f(x)$,

$$\begin{aligned} & \text{minimize} && f(x) \\ & \text{subject to} && 0 \leq x \leq c, \\ & && A^\top x = b, \quad b(V) = 0, \end{aligned} \quad (10)$$

the same descent applies with the arc cost replaced by the gradient: choose Δx as a negative cycle of the residual graph for the cost $\nabla f(x)$, so that $\nabla f(x)^\top \Delta x < 0$. The step is now limited by capacities *and* by the line search,

$$\alpha_{\text{cvx}} = \min\{\alpha_{\text{lin}}, t\},$$

where t is either an exact step, $t = \arg \min_{t>0} f(x + t \Delta x)$, or a backtracking step with parameters $\alpha \in (0, 1/2)$ and $\beta \in (0, 1)$ that shrinks t while

$$f(x + t \Delta x) > f(x) + \alpha t \nabla f(x)^\top \Delta x.$$

The linear problem (9) is the special case in which $t = +\infty$ and only the capacity bound is active. This convex cycle-cancelling descent is the method implemented by general-purpose network-optimization libraries such as LEMON.

5.4 Linear-fractional and statistical costs

Two quasi-convex costs arise repeatedly in timing. The *linear-fractional* min-cost flow,

$$\min \left\{ \frac{e^\top x + f}{g^\top x + h} : 0 \leq x \leq c, A^\top x = b, b(V) = 0 \right\},$$

is recast, with $p(x) = e^\top x + f$ and $q(x) = g^\top x + h$, as the family

$$\min \{ \gamma : (e - \gamma g)^\top x + (f - \gamma h) \leq 0, 0 \leq x \leq c, A^\top x = b \},$$

solved by bisection on γ with descent directions of cost $e - \gamma g$. This is the minimum cost-to-time ratio problem of Section 4 expressed as a flow: the affine parametric weight $d(\beta) = m - s\beta$ there is the affine cost $e - \gamma g$ here.

The *statistical* problem

$$\min \{ \Pr(\mathbf{d}^\top x > \alpha) : 0 \leq x \leq c, A^\top x = b, b(V) = 0 \}$$

treats $\mathbf{d}$ as a random vector with mean d and covariance Σ , so that $\mathbf{d}^\top x$ has mean $d^\top x$ and variance $x^\top \Sigma x$. The chance constraint is quasi-convex and is recast as

$$\min \{ \gamma : d^\top x + F^{-1}(1 - \gamma) \|\Sigma^{1/2} x\|_2 \leq \alpha, 0 \leq x \leq c, A^\top x = b \},$$

a second-order-cone (convex quadratic) constraint in x whose gradient is

$$d + F^{-1}(1 - \gamma) (\|\Sigma^{1/2} x\|_2)^{-1} \Sigma x.$$

This is the network analogue of the per-edge statistical constraints of Section 7: the covariance enters through the cone term, while the quantile $F^{-1}(1 - \gamma)$ plays the role of the yield parameter.

5.5 Additional constraints and the barrier method

A constraint that is not a network constraint—for instance $s^\top x \leq \gamma$ for a fixed vector s —breaks the network structure: the feasible set is no longer a network polytope, and (10) becomes a general convex program.

Such a constraint is handled by a logarithmic *barrier* [13],

$$\phi(x) = -\log(\gamma - s^\top x), \quad \nabla \phi(x) = \frac{s}{\gamma - s^\top x},$$

giving the approximation

$$\begin{aligned} & \text{minimize} && f(x) + \frac{1}{t} \phi(x) \\ & \text{subject to} && 0 \leq x \leq c, \quad A^\top x = b, \quad b(V) = 0, \end{aligned} \tag{11}$$

whose optimum tends to that of the original problem as $t \rightarrow \infty$. The *barrier method* alternates a *centering* step—minimizing $tf + \phi$, by Newton’s method in general—with an update $t \leftarrow \mu t$, $\mu > 1$, and stops when $1/t < \varepsilon$. The network-flow insight is that the centering step can itself be taken with a negative cycle: in place of the Newton direction one may use a negative cycle of the residual graph weighted by $\nabla(tf + \phi)$. The barrier method thus decomposes a constrained convex network program into a sequence of unconstrained cycle-cancellation steps.

6 Delay Padding

6.1 Motivation and formulation

When the TCG contains a negative cycle, no clock schedule can repair the circuit; the physical delays themselves must change. *Delay padding* inserts extra delay into selected data paths, usually by swapping a fast cell for a slower one. Its natural formulation adds a nonnegative padding vector p :

$$\text{find } p, u \quad \text{such that} \quad y \leq d + p, \quad Au = y, \quad p \geq 0. \tag{12}$$

If the objective is to minimize the total inserted delay $\sum_e p_e$, then (12) is the *dual* of a standard minimum-cost flow problem, solvable by the network simplex method. This observation is appealing but, in modern designs, *impractical*: delay insertion is realized by cell substitution, so minimizing $\sum p$ does not reflect the true cost, and—more importantly—there may be no valid position to insert delay along some paths.

6.2 The four structural configurations

A better strategy is to decide *where* delay may be inserted *before* deciding *how much* to insert. This is *path relationship analysis* (PRA). For a pair of

registers i and j , the maximum-delay (setup) path and the minimum-delay (hold) path interact in one of four ways, and each interaction determines how the TCG must be modified (Figure 6):

1. **Type (a), no insertion.** The two paths coincide or cannot be modified; no delay may be inserted independently, and the TCG is unchanged.
2. **Type (b), independent.** The setup and hold paths are disjoint; delays p_s and p_h may be inserted independently, modeled by auxiliary nodes connected to the original vertices.
3. **Type (c), shared.** The two paths overlap entirely, so a single delay must be shared: $p_s = p_h$.
4. **Type (d), containment.** The hold path is a sub-path of the setup path; padding upstream affects both, enforcing $p_s \geq p_h$.

6.3 Delay padding as a modified feasibility problem

With the TCG augmented by auxiliary nodes and coupling edges, the question “can the timing be fixed by padding?” again becomes a feasibility problem on a modified graph: an infeasible instance returns a negative cycle, which certifies that padding alone is insufficient and that D_{ij} must be reduced or T_{CP} increased. In the modified graph the padding is a *potential* on the auxiliary edges, so the whole machinery of Section 3 applies unchanged. This “determine the position first, then the value” strategy is both more flexible and more faithful to physical design than the minimize- $\sum p$ flow formulation.

6.4 Discrete padding and its complexity

Formulation (12) treats the padding p as a continuous, nonnegative real vector. In a standard-cell flow this is an idealization: delay is inserted either by swapping a cell for a slower one from a discrete library or by adding whole buffer stages, so the true feasible set of p is *discrete* (and, per path, non-convex), and the exact problem is a mixed-integer linear program (MILP) rather than a network flow. The continuous relaxation studied here should therefore be read as (i) a lower bound on the achievable clock period and (ii) a guide to *which* paths are worth padding. Path relationship analysis narrows the search by pre-selecting physically realizable positions and by coupling the setup and hold paddings ($p_s = p_h$ or $p_s \geq p_h$), but it does not make the values continuous. Closing the gap between a relaxed padding and a discrete, cell-level solution is a separate combinatorial problem; choosing positions

before values, as above, is what keeps that second problem tractable in practice.

Discreteness does not by itself place the problem beyond convex machinery. The level-set and cutting-plane methods that underlie the solvers are not restricted to continuous variables: some discrete problems whose relaxations are integral—minimum-cost matching being the classic example—are solved exactly by the same separation logic [9]. Extending (12) to a discrete technology library therefore appears feasible, although it requires real cell-level data; the continuous relaxation studied here is the first step.

7 Yield-Driven Scheduling under Process Variations

7.1 Timing yield

When process variations grow large enough, timing-failure-induced yield loss becomes a first-order concern, and the design objective shifts from the shortest nominal period to the *highest timing yield*. A circuit is called *functionally correct* for a given sample of process parameters if all of its setup- and hold-time constraints are satisfied; the *timing yield* is the fraction of manufactured samples that are correct [2], [14], [15], [16].

Maximizing yield is fundamentally different from minimizing a nominal period: it forces the schedule to reason about the *distributions* of the path delays rather than their worst-case values.

7.2 Primitive solutions and their shortcomings

Three elementary heuristics make the difficulty concrete.

1. **Margin pre-allocation.** Reserve a fixed timing margin Δd at both ends of every FSR,

$$\begin{aligned} \underline{w}_{ij} &\leq y_{ij} \leq \bar{w}_{ij} \\ \implies \underline{w}_{ij} + \Delta d &\leq y_{ij} \leq \bar{w}_{ij} - \Delta d \end{aligned}$$

and then optimize the clock period. The margin is usually the maximum timing uncertainty, which is too pessimistic, and a single fixed Δd cannot reflect the different uncertainties of different data paths.

2. **Least center error square (LCES).** Since a skew near the center of its FSR leaves the most slack on both sides, one may place each y_{ij} as close as possible to the midpoint of $[\underline{w}_{ij}, \bar{w}_{ij}]$.




Figure 6: The four structural configurations for delay insertion. Type (a): no feasible insertion; (b): p_s, p_h independent; (c): $p_s = p_h$; (d): $p_s \geq p_h$.




Figure 7: After period minimization, many skews sit on the boundary of their feasible skew regions; the narrow margins that remain are what process variations consume, which motivates yield-driven scheduling.

Writing $m_{ij} = \bar{w}_{ij} - \underline{w}_{ij}$ for the width of the FSR and using lower and upper fractions $\ell_k, u_k \in [0, 0.5]$,

$$\underline{w}_{ij} + \ell_k m_{ij} \leq y_{ij} \leq \bar{w}_{ij} - u_k m_{ij},$$

$$\min \sum_k \left(\frac{1}{2} - \min(\ell_k, u_k) \right)^2,$$

a small quadratic program [2], [14]. Minimizing the *total* center error, however, can drive individual slacks to zero, which is not optimal for yield.

3. **Incremental slack distribution.** Distribute slack greedily while checking all skew constraints [17]; the method again ignores the differences between path delays.

All three fail for the same reason: they do not weight each FSR by the *variance* of the underlying path delays.

7.3 Statistical timing model

Under process variations the maximum and minimum path delays become random variables $\tilde{D}_{ij}$ and $\tilde{d}_{ij}$, and the setup- and hold-time constraints become probabilistic:

$$T_{\text{skew}}(i, j) \leq T_{\text{CP}} - \tilde{D}_{ij} - T_{\text{setup}},$$

$$T_{\text{skew}}(i, j) \geq T_{\text{hold}} - \tilde{d}_{ij},$$

where $T_{\text{skew}}(i, j) = t_i - t_j$. After statistical static timing analysis (SSTA), every edge of the timing constraint graph carries a pair (μ_{ij}, σ_{ij}) —the mean and

standard deviation of its delay—so the schedule must bound the *probability* of violation rather than its worst case (Figure 8).




Figure 8: A statistical timing constraint graph: after SSTA each edge carries a pair (μ, σ) —mean and standard deviation—instead of a single deterministic weight.

7.4 EVEN: uniform slack maximization

The EVEN method assumes that all path delays have equal variance and maximizes a uniform margin β :

$$\max \beta \quad \text{s.t.} \quad t_j - t_i \leq \mu_{ij} - \beta.$$

This is the *minimum mean cycle* problem, whose optimum is the mean edge weight of the critical cycle C (the first negative cycle),

$$\beta^* = \frac{1}{|C|} \sum_{(i,j) \in C} \mu_{ij},$$

computable by Karp’s algorithm or by the faster methods of Dasdan and Gupta [18], [19].

EVEN then proceeds iteratively, much like minimum balancing: it identifies the most timing-critical cycle, distributes the slack evenly along it, freezes the corresponding skews, replaces the cycle by a super-vertex, and repeats (Figure 9). For the resulting schedule the slack of each edge is

$$\text{slack}_{ij} = T_{\text{CP}} - D_{ij} - T_{\text{setup}} - \text{skew}_{ij}.$$

EVEN’s weakness is exactly its assumption: the timing uncertainty of a long combinational path is normally larger than that of a short one, so an even slack distribution is not yield-optimal whenever the path delays along a critical cycle differ.

7.5 PROP: variance-proportional slack

PROP models each gate delay as Gaussian with a common variance and a path as the sum of n such gates, so a path delay is $\mathcal{N}(n\mu, n\sigma^2)$ and its standard deviation grows as the square root of the path length. It distributes the slack along the most critical cycle in proportion to those square roots by updating the TCG weights with a parameter α ,

$$\begin{aligned}\bar{w}_{ij}^{\text{PROP}} &= T_{\text{CP}} - (D_{ij} + \alpha\sqrt{D_{ij}}\sigma) - T_{\text{setup}}, \\ \underline{w}_{ij}^{\text{PROP}} &= -T_{\text{hold}} + (d_{ij} - \alpha\sqrt{d_{ij}}\sigma),\end{aligned}$$

where α ensures a minimum timing margin for each constraint. For a fixed clock period, α is increased until Bellman–Ford reports infeasibility; the largest feasible α gives every edge of the most critical cycle a slack equal to its pre-allocated margin, after which the remaining skews are filled in by EVEN. PROP’s weaknesses are its assumption of a common gate distribution and its unjustified use of the square root of the path delay as the weighting.

7.6 FP-PROP: false-path awareness

A *false path* can never be sensitized and therefore never carries a signal (Figure 10). If false paths are treated as real, non-critical cycles are mistaken for critical ones, slack is diverted to them, and the genuine critical cycles receive too little margin; the overall timing yield then drops. FP-PROP therefore performs a sensitizable-critical-path search and removes false paths from the scheduling problem [15].

7.7 Most critical cycle and C-PROP

Traditionally the most critical cycle is the *minimum mean cycle*, weighted by $\sum \mu_{ij}/|C|$. Under unequal variances the sound criterion weights each edge by its own uncertainty, and the critical cycle is the one maximizing the ratio $\sum \mu_{ij}/\sum \sigma_{ij}$. C-PROP is exactly the corresponding slack-maximization problem,

$$\max \beta \quad \text{s.t.} \quad t_j - t_i \leq \mu_{ij} - \sigma_{ij}\beta,$$

a *minimum cost-to-time ratio cycle* problem whose optimum is

$$\beta^* = \frac{\sum_{(i,j) \in C} \mu_{ij}}{\sum_{(i,j) \in C} \sigma_{ij}},$$

solved efficiently by Howard’s algorithm [17]. C-PROP reduces to EVEN when all σ_{ij} are equal, and as a variance tends to zero it assigns only a minimal margin to that constraint while leaving more for the others—precisely the desired behavior.

7.8 Whole flow: super-vertex contraction

Once the arrival times on the most critical cycle are fixed, the cycle is replaced by a super-vertex v' . An in-edge (u, v) from an outside vertex u to a cycle member v becomes an in-edge (u, v') of mean $\mu(u, v) - T_v$, and an out-edge (v, u) becomes (v', u) of mean $\mu(v, u) + T_v$; the variance of the edge weight is unchanged, and parallel edges may remain. The process repeats until the graph is reduced to a single super-vertex or has no edges left. The final arrival time of a register is then the sum of the frozen arrival times along its path in the contraction tree, $T_i^{\text{final}} = \sum_v T_v$ (Figure 12).

7.9 One parametric family

Minimum-period scheduling, EVEN, and C-PROP are three members of the single family (7), differing only in the objective $g(\beta)$ and the right-hand sides $f_{ij}(\beta)$ (Table 1).

Here β plays the role of the clock period in the Min. CP row and of the uniform margin in the EVEN and C-PROP rows. The objective g and the bounds f_{ij} need not be linear: any *monotone decreasing* functions will do, and the problem then still has a unique solution.

7.10 Yield maximization and the correlation question

The statistical objectives above are instances of the max-min formulation

$$\max \left\{ \min_{ij} \Pr\{t_j - t_i \leq \tilde{W}_{ij}\} \right\},$$

where $\tilde{W}_{ij}$ is the random edge slack. This is *not* exactly the timing yield, but it is reasonable, easy to solve, and—crucially—it needs no correlation information among the $\tilde{W}_{ij}$. Writing F_{ij} for the CDF of $\tilde{W}_{ij}$, it is equivalent to

$$\begin{aligned}\max \beta \quad \text{subject to} \quad & t_i - t_j \leq T_{\text{CP}} - F_{ji}^{-1}(\beta), \\ & t_j - t_i \leq F_{ij}^{-1}(1 - \beta).\end{aligned}$$

Because every CDF is monotone increasing, its inverse is well defined and the problem remains a parametric potential problem.




Figure 9: The iterative EVEN (minimum-balancing) process on a small timing constraint graph: identify the most critical cycle, distribute its slack evenly, freeze the arrival times, contract the cycle to a super-vertex, and repeat until a single vertex remains.

Table 1: Clock skew scheduling as one parametric family. Rows give the objective $g(\beta)$ and the setup and hold edge bounds; every constraint reads $t_i - t_j \leq f_e(\beta)$.

Problem	$g(\beta)$	$f_{ij}(\beta)$ (setup)	$f_{ji}(\beta)$ (hold)
Min. CP	$-\beta$	$\beta - D_{ij} - T_{\text{setup}}$	$d_{ij} - T_{\text{hold}}$
EVEN	β	$T_{\text{CP}} - D_{ij} - T_{\text{setup}} - \beta$	$d_{ij} - T_{\text{hold}} - \beta$
C-PROP	β	$T_{\text{CP}} - D_{ij} - T_{\text{setup}} - \sigma_{ij}\beta$	$d_{ij} - T_{\text{hold}} - \sigma_{ij}\beta$




Figure 10: A false path: although a structural path exists, no input pattern sensitizes it, so it carries no signal and must be excluded from the timing constraints.




Figure 11: The optimum of the minimum cost-to-time ratio cycle for a C-PROP example.




Figure 12: The contraction tree: each super-vertex merges a critical cycle, and the final arrival time of a register is the sum of the frozen arrival times along its root path—for example, $T_1^{\text{final}} = T_1 + T_7 + T_9$.

The reason no correlation is required is that the objective bounds each *individual* constraint separately, so only the *marginal* distributions are needed. This is not an assumption that the delays are uncorrelated: correlation is accounted for upstream, in the STA/SSTA step that produces $\tilde{D}_{ij}$ and $\tilde{d}_{ij}$, and once those marginals are fixed the per-edge objective needs no further correlation information. The price is that the result is a *surrogate* for, not the true, timing yield, which is a property of the *joint* distribution of the path delays. Replacing the surrogate by the true joint yield would require the full correlation matrix, which contemporary SSTA does not report, and remains an open problem (Section 11). Figure 13 compares the schedules produced by the methods above.




Figure 13: A comparison of the yield-driven scheduling methods on a common example.

7.11 Gaussian model and linearization

When the edge slacks are Gaussian, the probabilistic constraints can be written explicitly. With $z = \sqrt{2} \text{erf}^{-1}(2\beta - 1)$, maximizing the yield parameter β becomes

$$\begin{aligned} & \text{maximize} && \beta \\ & \text{subject to} && t_i - t_j \leq T_{\text{CP}} - \mu_{ij}^D - \sigma_{ij}^D z, \\ & && t_j - t_i \leq \mu_{ij}^H - \sigma_{ij}^H z. \end{aligned}$$

Since erf^{-1} is anti-symmetric and monotone, the substitution $\beta' = z$ linearizes the constraints into a minimum cost-to-time ratio problem,

$$t_i - t_j \leq T_{\text{CP}} - \mu_{ij}^D - \sigma_{ij}^D \beta', \quad t_j - t_i \leq \mu_{ij}^H - \sigma_{ij}^H \beta',$$

solvable by binary search on β' .

7.12 Log-normal model

A log-normal model replaces the Gaussian quantile by an exponential one, giving

$$t_i - t_j \leq T_{\text{CP}} - \exp(\mu_{ij}^D + \sigma_{ij}^D \beta'),$$

$$t_j - t_i \leq \exp(\mu_{ij}^H - \sigma_{ij}^H \beta').$$

This bypasses the explicit evaluation of the error function. The problem is nonlinear and non-convex, but it is still monotone in β' and can be solved efficiently by binary search.

7.13 Experimental results

Figure 14 reproduces the lecture’s experimental results for the yield-driven scheduling methods.

Circuit	CP	Even		Prop		In-Prop	
		Yield (%)	Yield (%)	Imp. (%)	Yield (%)	Imp. (%)	
s1423	55.73	74.1	75.5	1.9	75.5	1.9	
s1488	16.62	72.3	75.4	4.3	NA	NA	
s1494	16.62	77.9	78.8	1.2	NA	NA	
s5378	22.50	63.8	64.2	0.6	NA	NA	
s9234.1	40.86	74.1	83.9	13.2	NA	NA	
s13207.1	52.73	60.2	60.2	0.0	NA	NA	
s35932	31.96	64.3	98.6	53.3	98.6	53.3	
s38417	34.07	74.1	74.7	0.8	89.1	20.2	
s38584.1	50.24	72.5	85.8	18.3	NA	NA	
b04s	23.88	67.4	82.6	22.6	NA	NA	
b05s	47.68	67.2	86.9	29.3	86.9	29.3	
b06	4.92	63.1	75.1	19.0	NA	NA	
b07s	26.23	74.8	72.7	-2.8	72.7	-2.8	
b08	14.90	68.9	65.6	-4.8	65.6	-4.8	
b09	7.68	62.4	78.3	25.5	NA	NA	
b10	9.71	73.5	73.7	0.3	NA	NA	
b11s	29.86	71.5	58.6	-18.0	80.9	13.1	
b12	12.12	78.2	87.6	12.0	87.6	12.0	

Figure 14: Experimental results for the yield-driven scheduling methods.

7.14 Co-optimization of clock period and yield

When the clock period and the yield are optimized jointly, the objective is the ratio

$$\min \frac{T_{CP}}{\beta} \quad \text{subject to} \quad y_{ij} \leq T_{CP} - F_{ij}^{-1}(\beta) \quad (i, j) \in E_s,$$

The setup bound shrinks linearly with the period and grows with the yield, while the hold bound grows with the yield; both involve the quantile F_{ij}^{-1} , which is not concave over the whole interval $[0, 1]$. In practice only high-yield operation is of interest—say $\beta \geq 0.8$ —and on that range the problem is convex. Introducing the epigraph variable γ linearizes the ratio,

$$\min \gamma \quad \text{subject to} \quad T_{CP} - \gamma\beta \leq 0, \quad y_{ij} \leq T_{CP} - F_{ij}^{-1}(\beta), \quad y_{ji} \leq T_{ij}^{-1}(1-\beta), \quad T_{CP} \geq 0, \quad 0.8 \leq \beta \leq 1,$$

a convex network problem solved by binary search or the ellipsoid method, in the manner of Section 5.

7.15 Yield-driven delay padding and its dual

The padding formulation of Section 6 and the yield formulation of the present section can be merged into a single program that trades yield against the cost of inserted delay. With a padding vector $p \geq 0$, a weight γ on the yield, and per-edge Gaussian delays $\mathbf{d}_{ij}$ of mean d_{ij} and variance s_{ij} , the problem is

$$\max \gamma\beta - c^\top p \quad \text{subject to} \quad \beta \leq \Pr\{y_{ij} \leq \mathbf{d}_{ij} + p_{ij}\}, \quad Au = y, \quad p \geq 0$$

which is equivalent to the linear program

$$\max \gamma\beta - c^\top p \quad \text{subject to} \quad y \leq d - \beta s + p, \quad Au = y, \quad p \geq 0.$$

Its dual is a min-cost flow with one additional constraint,

$$\min d^\top x \quad \text{subject to} \quad 0 \leq x \leq c, \quad A^\top x = b, \quad b(V) = 0, \quad s^\top x \leq \gamma,$$

in which the non-network constraint $s^\top x \leq \gamma$ is exactly the one handled by the barrier method (11). Sections 6 and 7 are therefore primal and dual views of a single problem, which is why they share the same negative-cycle core.

8 Non-Gaussian and Heavy-Tailed Delay Models

8.1 Why the Gaussian model fails

$y_{ji} \leq F_{ij}^{-1}(1-\beta)$ ($(j, i) \in E_h$), $T_{CP} \geq 0$, $0 \leq \beta \leq 1$. At 65nm and below, the distribution of a *maximum* path delay is typically asymmetric and heavy-tailed, so the central limit theorem does not rescue the Gaussian assumption. Setup failures occur in slow corners and depend on the upper tail; hold failures occur in fast corners and depend on the lower tail. A model that fits the bulk but misses the tails therefore mis-estimates yield precisely where it matters (Figure 15).

8.2 Unimodality and quantile functions

A continuous distribution with mode m is *unimodal* if its CDF is convex for $x < m$ and concave for $x > m$. Normal, log-normal, and log-logistic distributions are unimodal, and their *quantile functions* are available in closed form:



Figure 15: Path-delay distributions are not Gaussian: they are asymmetric and heavy-tailed, so a Gaussian fit captures the bulk but understates the tails that govern setup and hold failures (illustration from the lecture).

$$\begin{aligned} \text{Normal:} & \quad \mu + \sigma\sqrt{2}\operatorname{erf}^{-1}(2p-1), \\ \text{Log-normal:} & \quad \exp(\mu + \sigma\sqrt{2}\operatorname{erf}^{-1}(2p-1)), \\ \text{Log-logistic:} & \quad \alpha\left(\frac{p}{1-p}\right)^{1/\kappa}. \end{aligned}$$

For the log-normal distribution the mode is $\exp(\mu - \sigma^2)$ and the CDF at the mode is $\frac{1}{2}(1 + \operatorname{erf}(-\sigma/\sqrt{2}))$.

8.3 Generalized extreme value distribution

The *generalized extreme value* (GEV) distribution [20] unifies the three extreme-value types and models skewness and tail weight through three parameters: location μ , scale $\sigma > 0$, and shape ξ . Its PDF and CDF are

$$\begin{aligned} f(x) &= \frac{1}{\sigma} t(x)^{\xi+1} e^{-t(x)}, & F(x) &= e^{-t(x)}, \\ t(x) &= \left[1 + \xi \frac{x - \mu}{\sigma}\right]^{-1/\xi} & (\xi \neq 0). \end{aligned}$$

The quantile function is

$$Q(\beta) = \mu + \frac{\sigma}{\xi} \left((-\ln \beta)^{-\xi} - 1 \right). \quad (13)$$

Because circuit delays are strictly positive and their maxima are heavy-tailed, the case $\xi \neq 0$ is the relevant one.

8.4 Statistical timing constraints with GEV

With GEV models for the maximum and minimum path delays, the setup and hold constraints of a critical path with skew T_{skew} become probabilistic,

$$T_{\text{skew}} \leq T_{\text{CP}} - \tilde{D} - T_{\text{setup}}, \quad T_{\text{skew}} \geq T_{\text{hold}} - \tilde{d},$$

and are converted to their deterministic equivalents through the quantile function (13):

$$\begin{aligned} T_{\text{CP}} - T_{\text{setup}} - T_{\text{skew}} &\geq Q_D(\beta), \\ T_{\text{hold}} - T_{\text{skew}} &\geq Q_d(1 - \beta). \end{aligned}$$

The setup constraint is sensitive to the *upper* tail $Q_D(\beta)$ (slow corners), whereas the hold constraint is sensitive to the *lower* tail $Q_d(1 - \beta)$ (fast corners). Because every CDF is monotone increasing, Q is monotone, so the yield parameter β remains a single scalar and the parametric framework of Section 4 still applies. In general Lawler’s binary search solves the resulting problem; for the special GEV form one may bypass the explicit quantile by searching directly on a transformed parameter, exactly as in the Gaussian linearization of Section 7.

9 Algorithms

9.1 Overview

Table 2 compares the principal algorithms for the parametric potential problem (6). They share one primitive: the detection and cancellation of negative cycles. The parametric shortest-path algorithm of Young, Tarjan, and Orlin [21] attains the best worst-case bound among them.

9.2 Lawler’s binary search

Lawler’s method brackets β and halves the interval with a feasibility test (Algorithm 1) at each step. It is easy to implement and robust to nonlinear $d(\beta)$, at the cost of repeated full-graph scans. It converges locally in the sense that each test is global but the bracket shrinks slowly (Figure 16).

9.3 Howard’s policy iteration

Howard’s algorithm is the method of choice for the single-parameter problem. It maintains one chosen outgoing arc per vertex (a *policy*); the policy graph

Table 2: Algorithms for the parametric potential problem. Here $n = |V|$ and $m = |E|$; ϵ is the binary-search tolerance.

Method	Strategy	Cost / behavior
Bellman–Ford	single feasibility test	$O(nm)$
Karp	minimum mean cycle	$O(nm)$
Lawler	binary search + feasibility test	$O(nm \log(1/\epsilon))$
Howard	policy iteration / cycle cancellation	fast in practice, exact
Young–Tarjan–Orlin	parametric shortest path	$O(nm + n^2 \log n)$




Figure 16: Lawler’s binary-search algorithm: bisect the bracket $[\beta_{\min}, \beta_{\max}]$ and test the timing constraint graph for a negative cycle; the upper or lower bound is updated until the tolerance is met.

contains a cycle, whose current ratio gives an improved β ; the potentials are then recomputed with respect to the parameterized weights and the policy is updated. Each iteration produces a strictly better β , and the algorithm returns the optimum together with its critical cycle. Variants that interleave binary-search pivots converge even faster, and hybrid Lawler–Howard methods are common in practice (Figures 17 and 18).




Figure 17: Howard’s policy-iteration algorithm: a cycle of the current policy yields an improved β , after which a new policy is computed from the parameterized shortest paths.

9.4 Karp’s minimum mean cycle

For the equal-variance case, the minimum mean cycle admits a closed-form characterization [18], computable in $O(nm)$ time from shortest-path distances with a bounded number of edges. Karp’s algorithm is the canonical subroutine for EVEN and for the minimum-mean-cycle form of any scheduling instance. The broader family of minimum-ratio algorithms is surveyed by Dasdan and Gupta [19].



Figure 18: Hybrid and improved variants interleave binary-search pivots with policy iteration, accelerating convergence.

9.5 From schedule to a sequentially realizable clock tree

The scheduling algorithms return arrival times, but these must be *realized* by a physical clock tree. Here the network viewpoint again pays off: the minimum balancing algorithm not only produces a schedule but also a contraction tree that encodes the order of criticality (Figure 5). Section 10 discusses how to exploit this structure.

Clock-Tree Synthesis and Co-optimization

10.1 Realizing a skew schedule

A clock tree delivering prescribed arrival times is built by *deferred-merge embedding* (DME) with Elmore delay and buffer insertion. Constraining the tree to deliver a *bounded* skew is possible but significantly complicates the algorithm. The practical recommendation is to keep the scheduled skews as a budget: if the schedule is over-optimized, the tree becomes hard to build, and budgeting techniques must be applied. After construction, a detailed timing analysis (not merely Elmore delay) should confirm the arrival times.

Realizing a prescribed skew is not free. Deferred-merge embedding places merging points and inserts buffers to meet target delays, and the buffer count, wirelength, and via count all grow with the *spread* of the required arrival times. Large skews between physically distant registers are the expensive case: they force detours in the clock wiring, raise clock-tree *power*, consume routing resources, and can induce *congestion* that ripples back into signal routing. A schedule that is optimal on the abstract TCG may therefore be unattractive, or even unrealizable, once these costs are counted. This is the central caveat of the approach: the algorithms here optimize a *delay-only* objective, and a production flow must re-optimize with the physical clock tree, or budget the skews so that the tree stays cheap to build.

10.2 Target skew versus actual skew

What the scheduler produces is a statement about *target* skews, and it must not be confused with the skew the silicon actually exhibits. The two concepts are distinct:

- the *target skew* is the value the scheduler intends to realize—the deterministic potential difference $y_{ij}^* = u_i - u_j$ returned by the optimization; and

- the *actual skew* is the value the built clock tree delivers, which is a *random variable*, because the buffers and wires that produce it are subject to process variation.

Writing the realized arrival time at register i as $u_i + \delta_i$, where δ_i is the random clock-arrival perturbation, the actual skew of path $i \rightarrow j$ is

$$y_{ij} = (u_i + \delta_i) - (u_j + \delta_j) = y_{ij}^* + (\delta_i - \delta_j).$$

The target skew is thus a deterministic *design intent*—we schedule a meeting at 10:00, not at 10:00 $\pm$ 34 minutes—whereas the actual skew is a stochastic *outcome* whose spread is governed by the correlation of the perturbations δ_i and δ_j . Two consequences follow. First, the scheduled skews are best treated as a *budget* rather than an exact requirement, since a tree that chases an over-optimized target becomes expensive and may still miss it, as noted above. Second, yield cannot be judged from the target skews alone; it requires a model of the actual skew as a random variable, which is exactly what the statistical methods of Section 7 supply. The target/actual distinction is therefore the interface between scheduling and clock-tree synthesis: the scheduler proposes a plan, and the tree—together with the variation it carries—determines what is realized.

The direction of causation reinforces the distinction. The schedule is computed *before* the clock tree exists, so it cannot be defined in terms of the skew that tree will eventually produce; the scheduler emits a target, and CTS then tries to approach it. The uncertainty of the underlying delays is already carried by the maximum and minimum path delays $\bar{w}$ and $\underline{w}$ supplied by STA/SSTA, so the target vector itself must remain a single, deterministic value, much as a class is scheduled for 03:25 p.m. and not for a distribution of starting times. This is why a reviewer who asks “why is the arrival time not a random variable when the method is statistical?” has mistaken the deterministic *target* for the random *outcome*—and, likewise, why an arrival time attributed to CTS is an *actual* skew, not the quantity optimized in Sections 4–7.

10.3 Criticality-ordered topology and placement

A useful heuristic is to make the clock-tree topology follow the criticality hierarchy: registers that belong to the same contracted cycle should share a branch. Two benefits follow. First, a lower branch has smaller skew *variation*, so the most critical registers—placed deepest—enjoy the strongest correlation and the largest

cancellation of variation. Second, if the placer also follows the topology—placing registers of the same branch physically together—spatial correlation further reduces skew variation. Since contemporary statistical timing analysis does not report the correlation information needed for a first-principles optimization, aligning topology, placement, and criticality is a valuable, practical surrogate.

10.4 Co-optimization after the tree is built

After clock-tree synthesis the delay estimates are more accurate, raising the question of whether to re-schedule. Some methods relax the constraint to $1.2u_i - 0.8u_j \leq w_{ij}$, reflecting that a built tree cannot realize an exact skew; the resulting problem is no longer a pure network flow but remains a linear program. It is also possible to relax the fixed arrival time to an *interval*, which reduces clock-tree wirelength and makes the schedule robust to unavoidable process variation.

11 Discussion and Open Problems

11.1 What the network formulation buys

Framing clock skew scheduling as a parametric potential problem yields three concrete benefits. First, *efficiency*: the combinatorial algorithms are faster than generic linear programming, and the parametric shortest-path methods are faster still. Second, *certificates*: an infeasible instance returns the most critical cycle, which is directly actionable, and a feasible schedule comes with a criticality hierarchy that guides clock-tree synthesis. Third, *extensibility*: delay padding and the yield-driven variants are all obtained by modifying the graph rather than the algorithm.

11.2 Scope and limitations

This article is a survey, not a report of new experiments. Its claims are of two kinds: mathematical statements—the reductions to the parametric potential problem and the negative-cycle criterion—which are exact under the stated model; and practical claims (efficiency, robustness), which are inherited from the cited literature and are not re-validated here.

The model makes four idealizations that a production flow must confront. (i) *Physical realizability*: arrival times are free continuous potentials, whereas a real

clock tree incurs power, area, buffer, and congestion costs that grow with skew spread (Section 10). (ii) *Discrete delay*: padding is quantized by cell substitution or buffer insertion, making the exact problem a MILP, of which the continuous network-flow formulation is a relaxation (Section 6). (iii) *Per-edge statistics*: the yield formulations use per-edge marginals and are not the true, joint timing yield; spatial correlation is accounted for upstream, in STA/SSTA, rather than in the scheduler (Section 7). (iv) *Flatness*: the difference-constraint model assumes one synchronous graph, whereas hierarchical timing abstractions, clock-domain crossings, and dynamic clock gating break that assumption. The framework is thus best understood as the exact *algorithmic core* of a larger physical-optimization problem, to which the surrounding engineering steps add costs and constraints that the core itself does not capture.

The absence of benchmark experiments is a matter of cost and access rather than relevance. Few groups carry out end-to-end useful-skew studies, and the enabling data—a manufactured test chip together with the spatial-correlation extraction that accompanies it—costs millions, which places a controlled comparison beyond the reach of a single algorithmic study. The position taken here is that the concepts must be settled first: a simple modeling error, such as treating the arrival time as a distribution or assuming Gaussian path delays, squanders exactly the data such an experiment is meant to produce. The open problems below are therefore stated so that groups with measurement access can pursue them.

11.3 Common misconceptions

Because the topic sits between algorithms and physical design, a few misreadings recur in reviews; they are worth stating explicitly.

- **“The arrival time should be a distribution because the method is statistical.”** The arrival time optimized here is a deterministic *target*, set before the clock tree exists; the random quantity is the *actual* skew the tree delivers (Section 10). The statistics enter through the path-delay bounds, not through the target vector.
- **“The formulation assumes uncorrelated delays.”** It does not. It is a per-edge objective that consumes only marginal path-delay distributions; correlation is accounted for upstream, when STA/SSTA produces those distributions (Section 7).
- **“Clock-tree synthesis is blamed for an un-**

fixable design.” When a scheduled skew proves hard to realize, the root cause is usually the *placement*: registers positioned without regard to the redefined critical paths. No amount of clock-tree effort recovers a poor register placement (Section 10).

- **“The model ignores power, area, and routing congestion.”** Those costs are real but belong to the physical stage, which the scheduler precedes. The recommended practice is to treat the schedule as a budget and to re-optimize once the tree exists (Section 10), rather than to burden the network-potential core with costs it cannot represent.
- **“The flat-graph model cannot handle hierarchy.”** The abstraction is a boundary condition, not a claim about the design: whatever the timing engine can characterize—hierarchical or not—is supplied to the scheduler as path delays. The hard problem lies in STA/SSTA itself, not in the scheduling core (Section 11.2).
- **“There are no benchmark experiments.”** The omission reflects cost and access, not relevance; the empirical claims are inherited from the literature, and the open problems are stated for groups with measurement access (Section 11.2).

11.4 Open problems

Several questions remain open.

- **Correlation-aware yield.** The max-min formulations here do not require correlation information, but they are not the true timing yield. Incorporating spatial correlation, as captured by Gaussian processes, remains a challenge.
- **Quantized delay padding.** Delay insertion by cell swapping is inherently quantized; integer or mixed-integer versions of (12) deserve more attention.
- **Non-Gaussian, multivariate yield.** Fitting GEV marginals is only half the story; the joint distribution and its tails govern the true yield.
- **Co-optimization.** Joint placement, clock-tree synthesis, and skew scheduling under a single objective is still largely unsolved.
- **Physical cost of skew.** The formulations optimize arrival times without charging for clock-tree power, area, buffer count, or routing congestion; a cost-aware objective that prices large, physically distant skews is missing.
- **Hierarchical and gated designs.** Flat difference constraints assume a single synchronous graph; hierarchical timing, clock-domain crossings (CDCs), and dynamic clock gating call for de-

composed or interface-constrained formulations.

- **Independent evaluation.** The speedups quoted in this survey are those reported in the cited papers; an apples-to-apples study on modern multi-million-gate benchmarks remains to be done.

12 Conclusion

Clock skew scheduling turns the inevitable, unwanted skew of a clock network into a deliberate design variable. We have shown that the problem has a single mathematical core: the timing constraints are a system of difference constraints, feasibility is the absence of a negative cycle in the timing constraint graph, and the optimization variants—minimum clock period, maximum slack, maximum yield—are parametric shortest-path problems with fast combinatorial solvers. Building on this core, delay padding, yield-driven scheduling, and heavy-tailed GEV delay models extend the framework without changing its algorithmic heart. The critical cycle and the criticality hierarchy returned by these solvers double as actionable design guidance for clock-tree synthesis and placement. As technology scales and process variation dominates, these ideas—useful skew, negative-cycle analysis, and parametric optimization—will remain central to timing closure.

References

- [1] J. P. Fishburn, “Clock skew optimization,” *IEEE Transactions on Computers*, vol. 39, no. 7, pp. 945–951, 1990.
- [2] J. L. Neves and E. G. Friedman, “Optimal clock skew scheduling tolerant to process variations,” in *Proceedings of the 33rd annual design automation conference (DAC)*, 1996, pp. 623–628. doi: [10.1145/240518.240636](https://doi.org/10.1145/240518.240636).
- [3] I. S. Kourtev and E. G. Friedman, *Timing optimization through clock-skew scheduling*. Boston, MA: Springer, 2000.
- [4] R. Bellman, “On a routing problem,” *Quarterly of Applied Mathematics*, vol. 16, no. 1, pp. 87–90, 1958.
- [5] T. H. Cormen, C. E. Leiserson, R. L. Rivest, and C. Stein, *Introduction to algorithms*, 3rd ed. Cambridge, MA: MIT Press, 2009.
- [6] E. L. Lawler, *Combinatorial optimization: Networks and matroids*. New York: Holt, Rinehart; Winston, 1976.

- [7] R. A. Howard, *Dynamic programming and markov processes*. Cambridge, MA: MIT Press, 1960.
- [8] L. G. Khachiyan, "A polynomial algorithm in linear programming," *Soviet Mathematics Doklady*, vol. 20, pp. 191–194, 1979.
- [9] M. Grötschel, L. Lovász, and A. Schrijver, *Geometric algorithms and combinatorial optimization*. Berlin: Springer, 1988.
- [10] X. Zhou, W.-S. Luk, H. Zhou, F. Yang, C. Yan, and X. Zeng, "Multi-parameter clock skew scheduling," *Integration, the VLSI Journal*, vol. 48, pp. 129–137, 2015, doi: [10.1016/j.vlsi.2014.07.005](https://doi.org/10.1016/j.vlsi.2014.07.005).
- [11] C. Albrecht, B. Korte, J. Schietke, and J. Vygen, "Cycle time and slack optimization for VLSI-chips," in *Proceedings of the IEEE/ACM international conference on computer-aided design (ICCAD)*, 1999, pp. 232–238. doi: [10.1109/ICCAD.1999.810654](https://doi.org/10.1109/ICCAD.1999.810654).
- [12] R. K. Ahuja, T. L. Magnanti, and J. B. Orlin, *Network flows: Theory, algorithms, and applications*. Englewood Cliffs, NJ: Prentice Hall, 1993.
- [13] S. Boyd and L. Vandenberghe, *Convex optimization*. Cambridge, UK: Cambridge University Press, 2004.
- [14] I. S. Kourtev and E. G. Friedman, "Clock skew scheduling for improved reliability via quadratic programming," in *Proceedings of the IEEE/ACM international conference on computer-aided design (ICCAD)*, 1999, pp. 239–243. doi: [10.1109/ICCAD.1999.810655](https://doi.org/10.1109/ICCAD.1999.810655).
- [15] J.-L. Tsai, D. H. Baik, C. C.-P. Chen, and K. K. Saluja, "Yield-driven, false-path-aware clock skew scheduling," *IEEE Design & Test of Computers*, vol. 22, no. 3, pp. 214–222, 2005.
- [16] C. Visweswariah, K. Ravindran, K. Kalafala, S. G. Walker, and S. Narayan, "First-order incremental block-based statistical timing analysis," in *Proceedings of the 41st annual design automation conference (DAC)*, 2004, pp. 331–336.
- [17] X. Wei, Y. Cai, and X. Hong, "Clock skew scheduling under process variations," in *Proceedings of the 7th international symposium on quality electronic design (ISQED)*, 2006, pp. 237–242.
- [18] R. M. Karp, "A characterization of the minimum cycle mean in a digraph," *Discrete Mathematics*, vol. 23, no. 3, pp. 309–311, 1978.
- [19] A. Dasdan and R. K. Gupta, "Faster maximum and minimum mean cycle algorithms for system-performance analysis," *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems*, vol. 17, no. 10, pp. 889–899, 1998.
- [20] A. F. Jenkinson, "The frequency distribution of the annual maximum (or minimum) values of meteorological elements," *Quarterly Journal of the Royal Meteorological Society*, vol. 81, no. 348, pp. 158–171, 1955.
- [21] N. E. Young, R. E. Tarjan, and J. B. Orlin, "Faster parametric shortest path and minimum-balance algorithms," *Networks*, vol. 21, no. 2, pp. 205–221, 1991.